



Generalized Additive Models (GAMs) with R

Day 3: What are GAMs?

**An Online Course
Presented by Geoffrey S. Hubona**

Day 3 Agenda (1 of 2)



- What are GAMs?
- We Begin with Simple Univariate Additive Model.
- Representing Smooth Functions: Regression Splines
 - Polynomial Basis
- Representing Smooth Functions: Cubic Splines
 - Knot Locations
 - Cubic Spline Basis
 - **Example in R: Engine Wear as a Function of Capacity**
 - **Example Functions in R: Cubic Spline; Model Matrix; Select Knots and Fit Model**
- Controlling the Degree of Smoothing
 - Penalize “Wiggleness”
 - **Example Functions in R: Penalized Regression Spline Penalty Matrix; Fitting Penalized Regression Spline**

Day 3 Agenda (2 of 2)



- Cross Validation
 - Choosing a Smoothing Parameter
 - **Example Scripts in R: Search for GCV-optimal Smoothing Parameter; Plot Optimal Engine Wear Model**
- Two Term Additive Models
 - **Example Functions in R: Penalized Regression Spline Two Term Additive Model; Fit Model and Calculate Optimal GCV Score**
 - **New Example in R: Estimate 2-term Additive Model for 31 Felled Cherry Trees.**
- Generalized Additive Models (GAMs)
 - **Revised Example in R: Generalized Additive Model for 31 Felled Cherry Trees.**

What Are GAMs?

SIMON N. WOOD

Parametric Only:
Row (x_i); Vector (β)

Smooth
Functions

Isotropic
Smooth

- A **Generalized Additive Model** (GAM) is a generalized linear model with a linear predictor involving a sum of smooth functions of covariates.
- General model structure of a GAM:

$$g(\mu_i) = X_i^* \Theta + f_1(x_{1i}) + f_2(x_{2i}) + f_3(x_{3i}, x_{4i}) + \dots$$

where $\mu_i \equiv E(Y_i)$, $Y_i \sim$ some exponential family distribution. Y_i is a response variable, X_i^* is a row of the model matrix for any strictly parametric model components, Θ is the corresponding parameter vector, and the f_j are smooth functions of the covariates, x_k .

Representing a Univariate Smooth Function using Regression Splines

- Here is an additive model containing one smooth function of one covariate:
$$y_i = f(x_i) + \varepsilon_i, \quad (1)$$

where y_i is a response variable, x_i a covariate, f is a smooth function and the ε_i are i.i.d. $N(0, \sigma^2)$ random variables. Also, we will assume that the x_i lie in the interval $[0, 1]$.

- **What is the function f ? How is f represented?**
- If we represent the smooth function f in a way that equation (1) becomes a linear model, we can estimate f .
- We choose a **basis** which defines a (limiting) space of functions (**'basis functions'**) which can be treated as if completely known (even though f is unknown).

Representing a Univariate Smooth Function using Regression Splines

- Here is the original model containing one smooth function of one covariate: $y_i = f(x_i) + \varepsilon_i, \quad (1)$

Repeat equation (1)
From previous slide

- Have chosen some basis functions, let's say for example, a 4th order polynomial, which can be treated as completely known: if $b_j(x)$ is the j^{th} such basis function, then f is assumed to have this representation:

$$f(x) = \sum_{j=1}^q b_j(x) \beta_j \quad (2)$$

- For some unknown values of the parameters, β_j .

Representing a Univariate Smooth Function using Regression Splines

- Here is the original model containing one smooth function of one covariate:

$$y_i = f(x_i) + \varepsilon_i, \quad (1)$$

Repeat equations (1) and (2) from previous slide.

- Function f now is: $f(x) = \sum_{j=1}^q b_j(x)\beta_j$ (2)

- If we substitute (2) into (1) we have $y_i = \sum_{j=1}^q b_j(x_i)\beta_j + \varepsilon_i$ which is a linear model.

- What does it look like?***

- If we believe f to be a 4th order polynomial, so that the space of polynomials of order 4 and below contains f , then a **basis** for this space is:

$$b_1(x) = 1, b_2(x) = x, b_3(x) = x^2, b_4(x) = x^3, b_5(x) = x^4$$

Representing a Univariate Smooth Function using Regression Splines

- So this is the **basis** of the space of polynomials of order 4 and below which contain f

$$b_1(x) = 1, b_2(x) = x, b_3(x) = x^2, b_4(x) = x^3, b_5(x) = x^4$$

- Then the original univariate additive model (equation (1)) becomes the simple linear model which we can estimate:

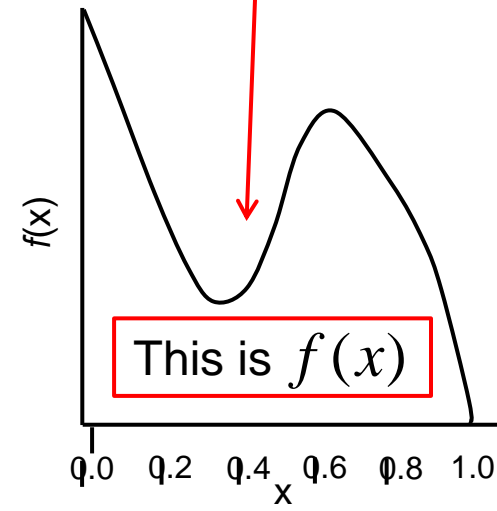
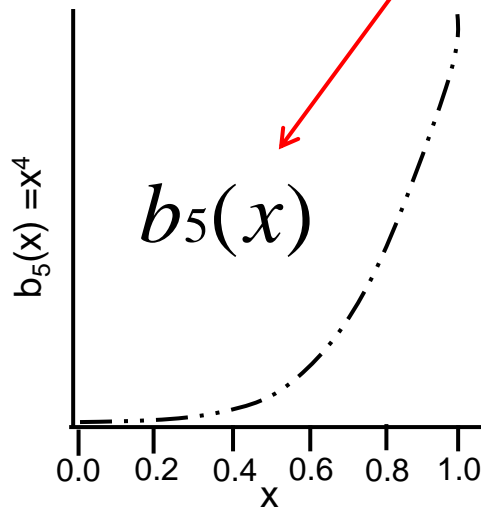
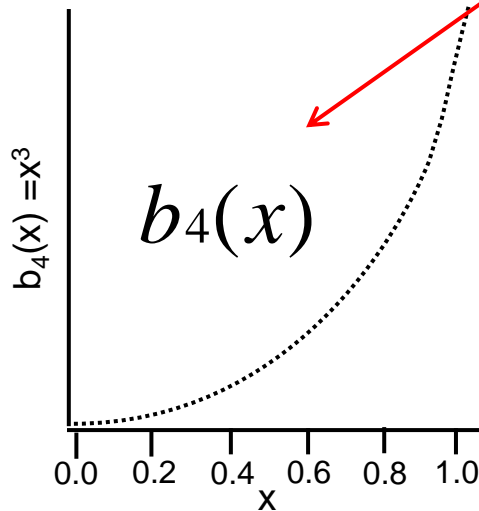
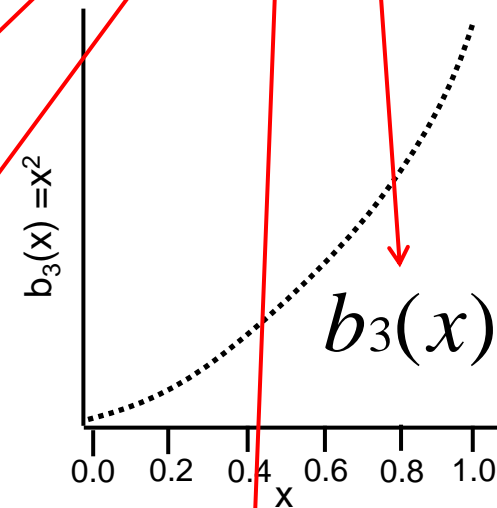
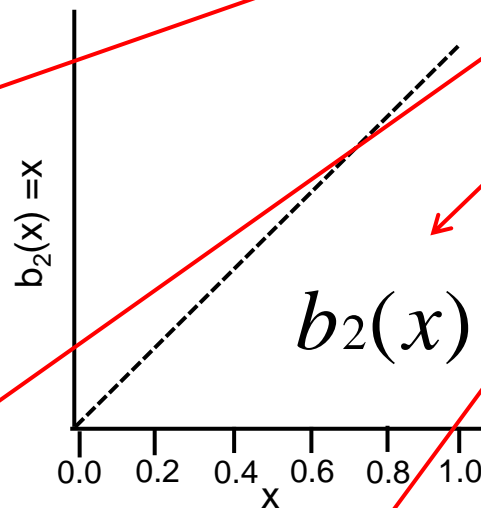
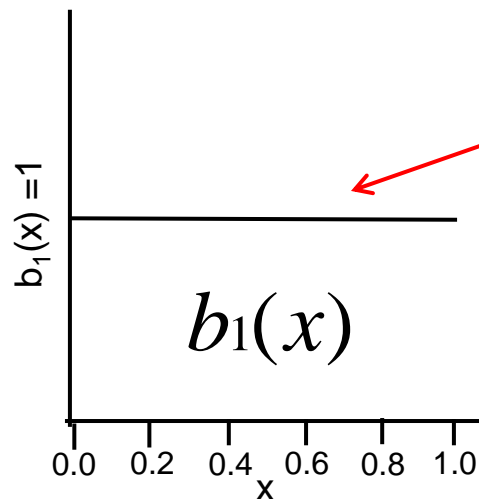
$$y_i = \beta_1 + x_i \beta_2 + x_i^2 \beta_3 + x_i^3 \beta_4 + x_i^4 \beta_5 + \varepsilon_i$$

- So for this simple univariate additive (not generalized) model, we need to represent the red circled term in terms of a **model matrix**:

$$y_i = \sum_{j=1}^q b_j(x_i) \beta_j + \varepsilon_i$$

Illustration of f as 4th Order Polynomial

Multiply each basis function by real valued parameter, β_j and then sum to give final curve.

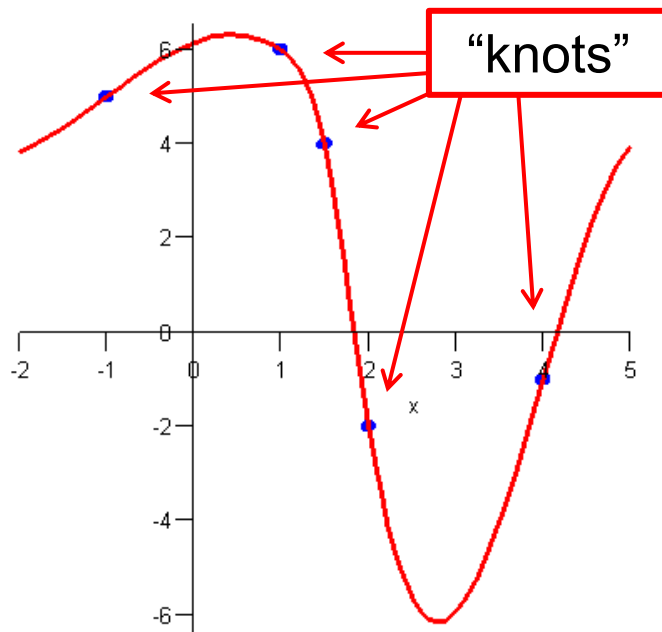


Note we have scaled x to $[0,1]$.

Cubic Splines



- One can represent a univariate function (eq (1) on slides #4-6) with a ***cubic spline***:
 - Is a curve made up of sections of a cubic polynomial joined together so they are continuous in value.



Conventional cubic spline:
knots occur where have data.
Regression splines: must
choose knot locations.

We denote **knot locations** by:

$$\{x_i^* : i = 1, \dots, q - 2\}$$

Question: Given knot locations, how do we establish the basis?

Knot Locations



- Given knot locations, there are alternative, but equivalent, ways of writing down a basis for cubic splines.
 - Knot locations are denoted by $\{x_i^* : i = 1, \dots, q - 2\}$
 - A simple, general basis function to use (Gu 2002) is:
 $b_1(x) = 1, b_2(x) = x$ and $b_{i+2} = R(x, x_i^*)$ for $i = 1 \dots q - 2$ where
- $$R(x, z) = \left((z - 1/2)^2 - 1/12 \right) \left((x - 1/2)^2 - 1/12 \right) / 4$$
$$- \left((|x - z| - 1/2)^4 - 1/2 (|x - z| - 1/2)^2 + 7/240 \right) / 24$$
- Using this cubic spline basis for f means that eq (1) becomes a linear model and so the model can be estimated by least squares.
 - We illustrate a rank 5 example of this basis two slides back.

Cubic Spline Basis



- Using this cubic spline basis for f means that eq (1) becomes a linear model: $y = X\beta + \varepsilon$
 - Where the i^{th} row of the model matrix is:
$$X_i = [1, x_i, R(x_i, x_1^*), R(x_i, x_2^*), \dots, R(x_i, x_{q-2}^*)]$$
- Steps: Cubic Spline Basis to Estimate a Model and Predict:
 - 1) Read in data, scale x axis $[0,1]$, plot data for spline fitting.
 - 2) Write a basis function (**rk()**)
 - 3) Write a function (**sp1.X()**) which will accept an array of x values and a sequence of knots to produce a model matrix for the spline.
 - 4) Select a set of knots.
 - 5) Fit the model.
 - 6) Generate x values for prediction.
 - 7) Create prediction matrix.
 - 8) Plot the fitted spline.

Example of Cubic Spline Basis in R

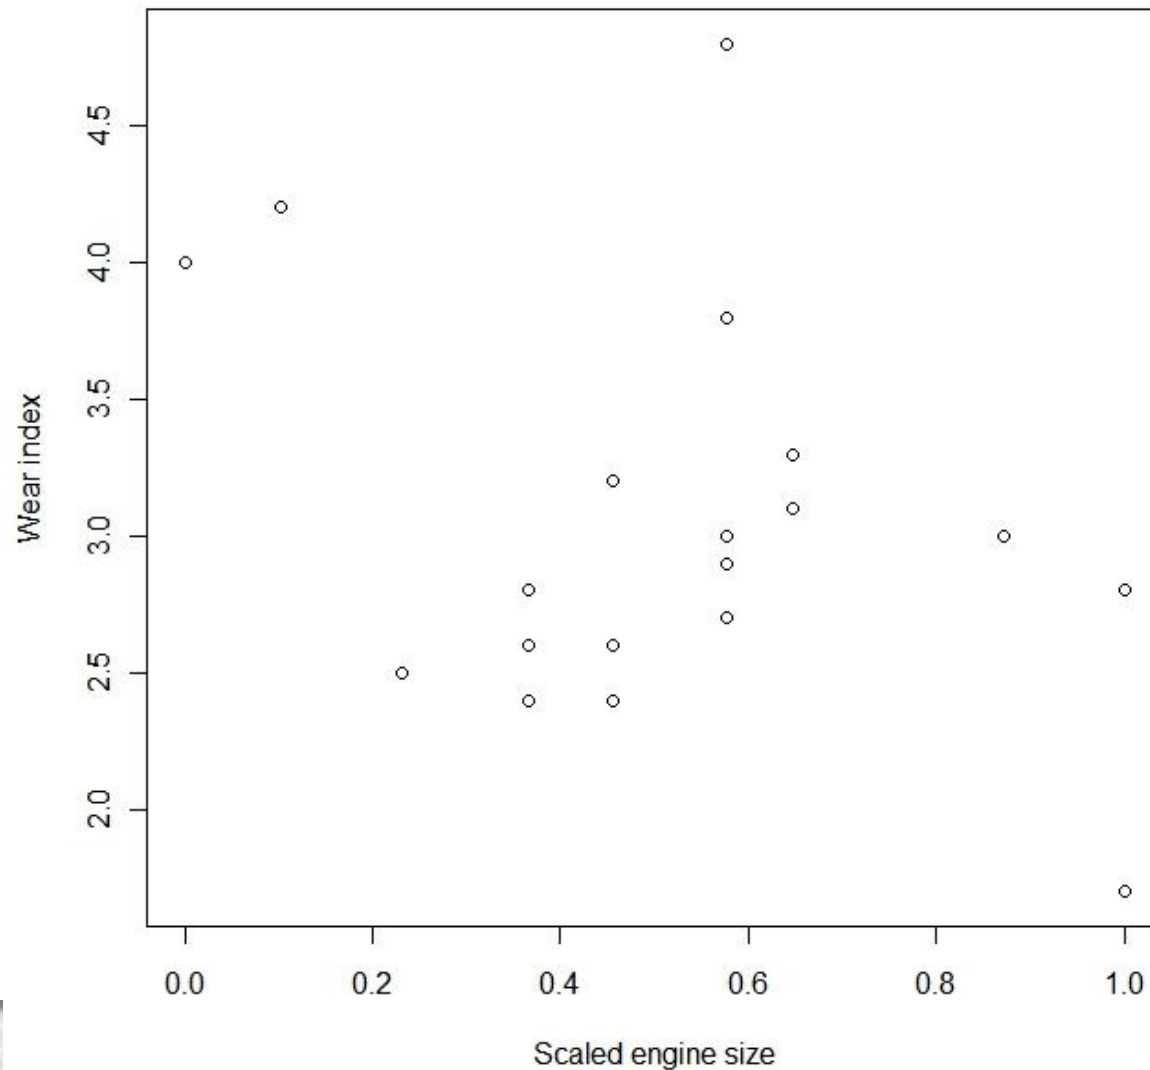


- People who do not know cars believe that a car engine with a larger cylinder capacity wears out less quickly than a car with a small engine capacity.
- We read the data into R and scale the engine capacity data to lie in $[0, 1]$ interval:

```
#### Using the cubic spline basis
size <- c(1.42,1.58,1.78,1.99,1.99,1.99,2.13,2.13,2.13,
          2.32,2.32,2.32,2.32,2.32,2.43,2.43,2.78,2.98,2.98)
wear <- c(4.0,4.2,2.5,2.6,2.8,2.4,3.2,2.4,2.6,4.8,2.9,
          3.8,3.0,2.7,3.1,3.3,3.0,2.8,1.7)
x <- size-min(size); x <- x/max(x)
plot(x,wear,xlab="Scaled engine size",ylab="Wear index")
```

This line scales the engine data

Plot of Wear Index by Scaled Engine Size



Function Defining Basis $R(x, z)$

- So now we write a function `rk()` defining the cubic spline $R(x, z)$ on the interval $[0, 1]$.

```
rk<-function(x,z) # R(x,z) for cubic spline on [0,1]
{ ((z-0.5)^2-1/12)*((x-0.5)^2-1/12)/4-
  ((abs(x-z)-0.5)^4-(abs(x-z)-0.5)^2/2+7/240)/24
}
```

From Gu (2002):

$$R(x, z) = \left((z - 1/2)^2 - 1/12 \right) \left((x - 1/2)^2 - 1/12 \right) / 4 \\ - \left((|x - z| - 1/2)^4 - 1/2 (|x - z| - 1/2)^2 + 7/240 \right) / 24$$

Function Defining Basis $R(x, z)$

- For this basis (a rank 6 basis):
- $b_1(x) = 1, b_2(x) = x$ and $b_{i+2} = R(x, x_i^*)$ for $i = 1, \dots, q-2$
- Where $R(x, z) = \left((z - 1/2)^2 - 1/12 \right) \left((x - 1/2)^2 - 1/12 \right) / 4$
$$- \left((|x - z| - 1/2)^4 - 1/2(|x - z| - 1/2)^2 + 7/240 \right) / 24$$
- Using this cubic spline basis for f means that the univariate smooth function $y_i = f(x_{1i}) + \varepsilon_i$, becomes a linear model $y = X\beta + \varepsilon$, where X is the model matrix.
- The i^{th} row of the model matrix is

$$Xi = \left[1, x_i, R(x_i, x_1^*), R(x_i, x_2^*), \dots, R(x_i, x_{q-2}^*) \right]$$

Function to Produce the Model Matrix for Spline



- We will use `rk()` function in the `spl.X()` function below to take a sequence of knots and an array of x values to produce a model matrix for the spline:

```
spl.X<-function(x,xk)
# set up model matrix for cubic penalized regression
# spline
{ q<-length(xk)+2    # number of parameters
  n<-length(x)       # number of data
  X<-matrix(1,n,q)   # initialized model matrix
  X[,2]<-x            # set second column to x
  X[,3:q]<-outer(x,xk,FUN=rk) # and remaining to R(x,xk)
  X
}
```

This is the 'model matrix'

Select Knots and Fit Model



- Select a set of knots, x_i^* , and the model can be fitted.
- We use rank 6 basis, so that $q = 6$ and there are 4 knots evenly spread over $[0,1]$:

```
xk<-1:4/5          # choose some knots
X<-spl.X(x,xk)     # generate model matrix
mod.1<-lm(wear~X-1) # fit model
xp<-0:100/100      # x values for prediction
Xp<-spl.X(xp,xk)   # prediction matrix
lines(xp,Xp%*%coef(mod.1)) # plot fitted spline
```

- The model fit (next slide) seems plausible, but the degree of model smoothness, controlled with the basis dimension (i.e. q , the number of knots + 2), was essentially arbitrary.

Model Matrix X Looks Like

Is a rank 6 cubic spline

```
> X
```

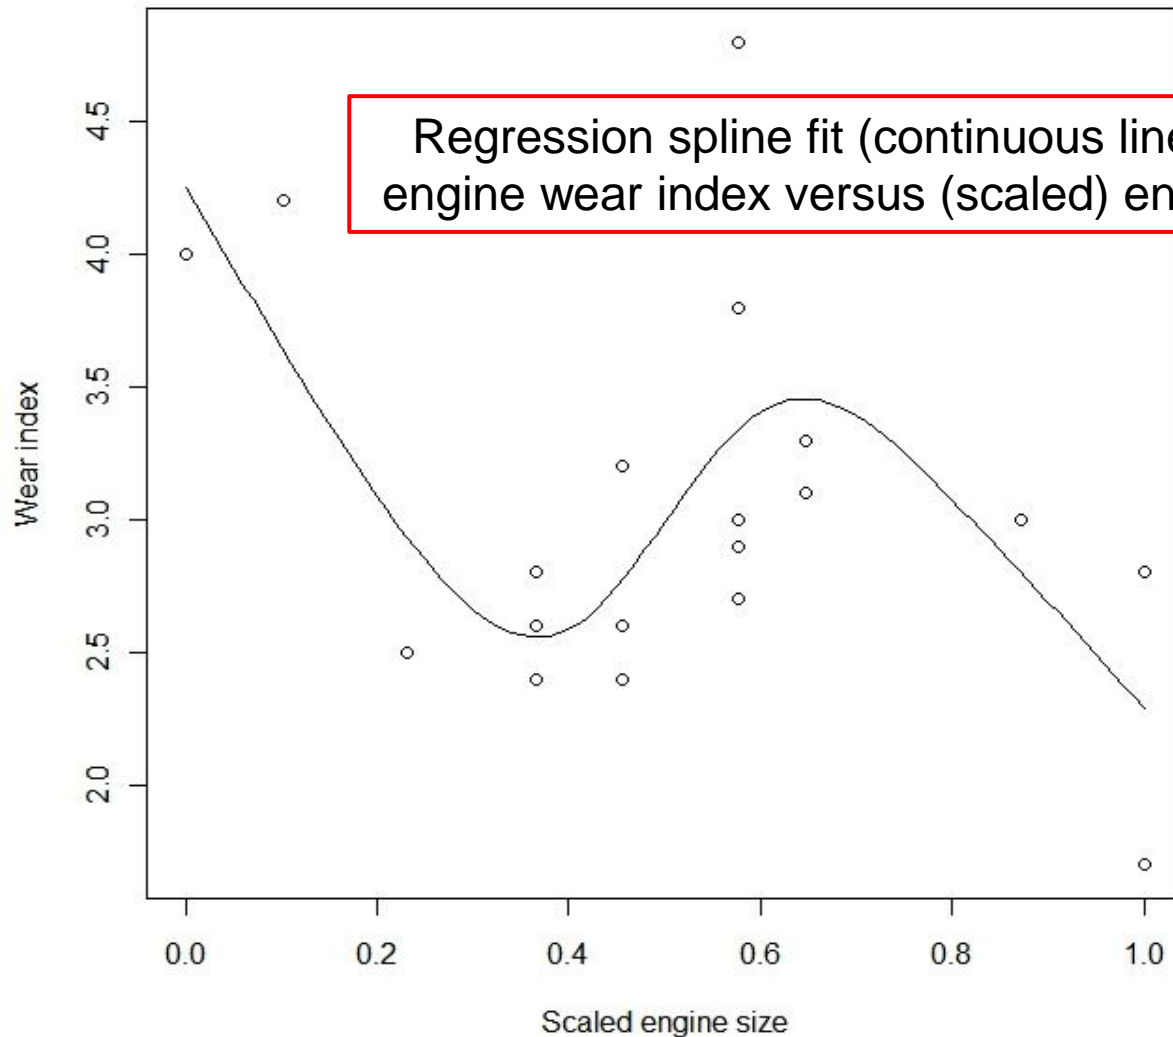
	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	0.0000000	0.0006000000	-0.0040666667	-0.0040666667	6.000000e-04
[2,]	1	0.1025641	0.0011910163	-0.0017986610	-0.0025832101	-3.421181e-04
[3,]	1	0.2307692	0.0013337512	0.0007641866	-0.0006723177	-1.134463e-03
[4,]	1	0.3653846	0.0004863279	0.0025379134	0.0012408700	-1.235660e-03
[5,]	1	0.3653846	0.0004863279	0.0025379134	0.0012408700	-1.235660e-03
[6,]	1	0.3653846	0.0004863279	0.0025379134	0.0012408700	-1.235660e-03
[7,]	1	0.4551282	-0.0002514095	0.0027666999	0.0022402848	-8.735914e-04
[8,]	1	0.4551282	-0.0002514095	0.0027666999	0.0022402848	-8.735914e-04
[9,]	1	0.4551282	-0.0002514095	0.0027666999	0.0022402848	-8.735914e-04
[10,]	1	0.5769231	-0.0010382837	0.0019246203	0.0027870084	8.294965e-06
[11,]	1	0.5769231	-0.0010382837	0.0019246203	0.0027870084	8.294965e-06
[12,]	1	0.5769231	-0.0010382837	0.0019246203	0.0027870084	8.294965e-06
[13,]	1	0.5769231	-0.0010382837	0.0019246203	0.0027870084	8.294965e-06
[14,]	1	0.5769231	-0.0010382837	0.0019246203	0.0027870084	8.294965e-06
[15,]	1	0.6474359	-0.0012606937	0.0010733705	0.0024330758	5.897512e-04
[16,]	1	0.6474359	-0.0012606937	0.0010733705	0.0024330758	5.897512e-04
[17,]	1	0.8717949	-0.0005452104	-0.0022051957	-0.0012497963	1.295347e-03
[18,]	1	1.0000000	0.0006000000	-0.0040666667	-0.0040666667	6.000000e-04
[19,]	1	1.0000000	0.0006000000	-0.0040666667	-0.0040666667	6.000000e-04

Prediction Model Matrix Xp

> Xp

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	0.00	6.000000e-04	-4.066667e-03	-4.066667e-03	6.000000e-04
[2,]	1	0.01	6.632829e-04	-3.843350e-03	-3.923317e-03	5.033829e-04
[3,]	1	0.02	7.262600e-04	-3.620140e-03	-3.779873e-03	4.070600e-04
[4,]	1	0.03	7.886162e-04	-3.397150e-03	-3.636250e-03	3.113163e-04
[5,]	1	0.04	8.500267e-04	-3.174507e-03	-3.492373e-03	2.164267e-04
[6,]	1	0.05	9.101562e-04	-2.952344e-03	-3.348177e-03	1.226563e-04
[7,]	1	0.06	9.686600e-04	-2.730807e-03	-3.203607e-03	3.026000e-05
[8,]	1	0.07	1.025183e-03	-2.510050e-03	-3.058617e-03	-6.051708e-05
[9,]	1	0.08	1.079360e-03	-2.290240e-03	-2.913173e-03	-1.494400e-04
[10,]	1	0.09	1.130816e-03	-2.071550e-03	-2.767250e-03	-2.362837e-04
[11,]	1	0.10	1.179167e-03	-1.854167e-03	-2.620833e-03	-3.208333e-04
[12,]	1	0.11	1.224016e-03	-1.638284e-03	-2.473917e-03	-4.028837e-04
[13,]	1	0.12	1.264960e-03	-1.424107e-03	-2.326507e-03	-4.822400e-04
[14,]	1	0.13	1.301583e-03	-1.211850e-03	-2.178617e-03	-5.587171e-04
[15,]	1	0.14	1.333460e-03	-1.001740e-03	-2.030273e-03	-6.321400e-04
[16,]	1	0.15	1.360156e-03	-7.940104e-04	-1.881510e-03	-7.023437e-04
[17,]	1	0.16	1.381227e-03	-5.889067e-04	-1.732373e-03	-7.691733e-04
[18,]	1	0.17	1.396216e-03	-3.866837e-04	-1.582917e-03	-8.324837e-04
[19,]	1	0.18	1.404660e-03	-1.876067e-04	-1.433207e-03	-8.921400e-04
[20,]	1	0.19	1.406083e-03	8.049583e-06	-1.283317e-03	-9.480171e-04
[21,]	1	0.20	1.400000e-03	2.000000e-04	-1.133333e-03	-1.000000e-03
[22,]	1	0.21	1.386083e-03	3.879496e-04	-9.833504e-04	-1.047984e-03
[23,]	1	0.22	1.364660e-03	5.715933e-04	-8.334733e-04	-1.091873e-03
[24,]	1	0.23	1.336216e-03	7.506162e-04	-6.838171e-04	-1.131584e-03

Regression Spline Fit



Regression spline fit (continuous line) to data on engine wear index versus (scaled) engine capacity.

Controlling the Degree of Smoothing



- So how do we ***control the degree of smoothing*** ?
- One obvious choice would be manipulate q via the number of knots (then backfit q and test fit hypotheses)
- Another: keep the basis dimension fixed, at a size slightly larger than believed necessary, but add a “wiggleness” penalty.
 - Instead of fitting the model by minimizing the sum of the least squares, minimize the sum of the least squares PLUS $\lambda \int_0^1 [f''(x)]^2 dx$, (**lambda times integrated square of second derivative**).
- ***Smoothing parameter*** λ controls trade off between model fit and model smoothness.
 - High λ leads to a straight-line estimate
 - $\lambda = 0$ results in an un-penalized regression spline estimate.

Penalized Regression Spline

Penalty Matrix

- Need to derive the square root of the penalized regression spline matrix **S**:

```
spl.S<-function(xk)
# set up the penalized regression spline penalty matrix,
# given knot sequence xk
{ q<-length(xk)+2;S<-matrix(0,q,q) # initialize matrix to 0
  S[3:q,3:q]<-outer(xk,xk,FUN=rk)  # fill in non-zero part
  S
}

mat.sqrt<-function(S) # A simple matrix square root
{ d<-eigen(S,symmetric=TRUE)
  rS<-d$vectors%*%diag(d$values^0.5)%*%t(d$vectors)
}
```

Function for Fitting Penalized Regression Spline

- To use this function, we need to choose the basis dimension q , the knot locations, x_j^* , and a value for the smoothing parameter, λ :

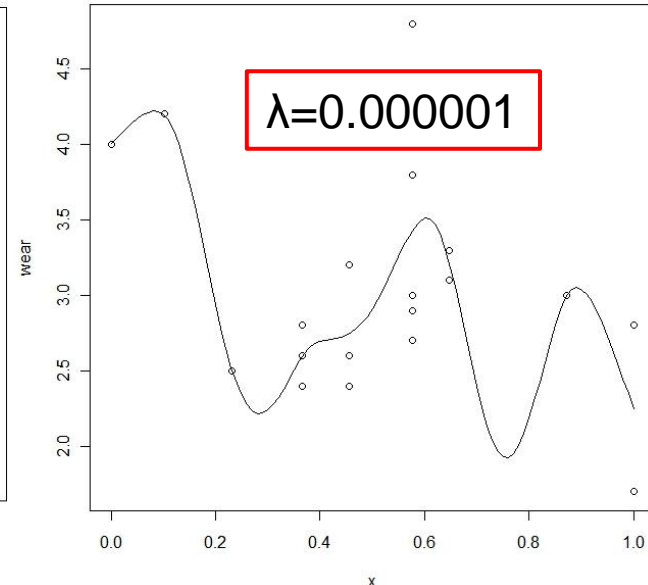
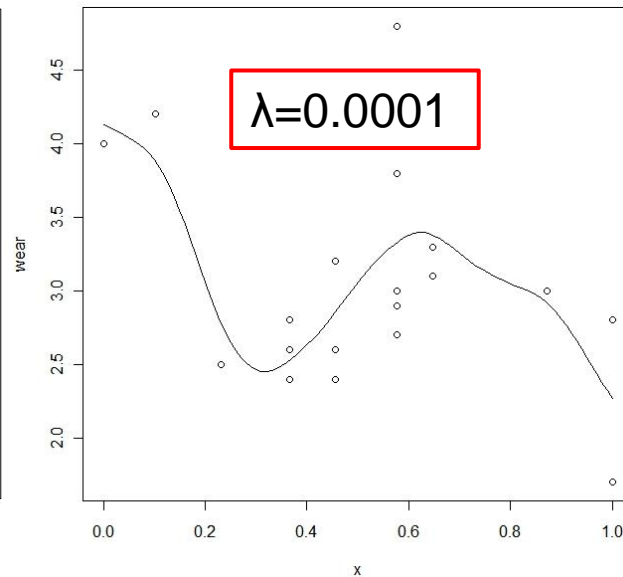
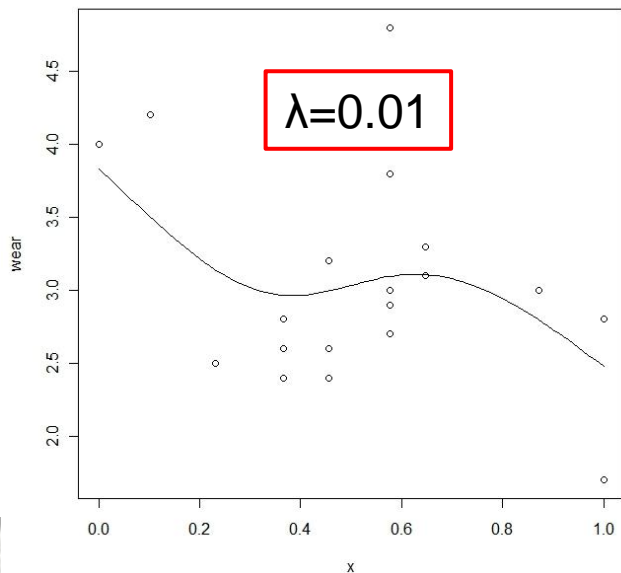
```
prs.fit<-function(y,x,xk,lambda)
# function to fit penalized regression spline to x,y data,
# with knots xk, given smoothing parameter, lambda.
{ q<-length(xk)+2          # dimension of basis
  n<-length(x)             # number of data
  # create augmented model matrix ....
  Xa <- rbind(spl.X(x,xk),mat.sqrt(spl.S(xk))*sqrt(lambda))
  y[(n+1):(n+q)]<-0      # augment the data vector
  lm(y~Xa-1) # fit and return penalized regression spline
}
```

Smoothing Parameter Controls the Smooth



- In this example, $q = 9$ and the knots are evenly spread over $[0,1]$. The smoothing parameter, $\lambda = 10^{-4}$, really controls the behavior of the fitted model:

```
xk<-1:7/8 # choose some knots  
mod.2<-prs.fit(wear,x,xk,0.0001) # fit pen. reg. spline  
Xp<-spl.X(xp,xk) # matrix to map params to fitted values at xp  
plot(x,wear);lines(xp,Xp%*%coef(mod.2)) # plot data & spl. fit
```



Choosing a Smoothing Parameter λ : Cross Validation

- If λ is too high, the data is over smoothed, if too low, the data is undersmoothed.
- In both cases, the spline estimate \hat{f} will not be close to the true function f : This is our goal.
- Imagine the criterion **choose λ to minimize**:

$$M = \frac{1}{n} \sum_{i=1}^n (\hat{f}_i - f_i)^2$$

- Since f is unknown, M cannot be used directly but you can derive an estimate of $E(M) + \sigma^2$ which is the expected square error in predicting a new variable.

Choosing a Smoothing Parameter λ : Cross Validation

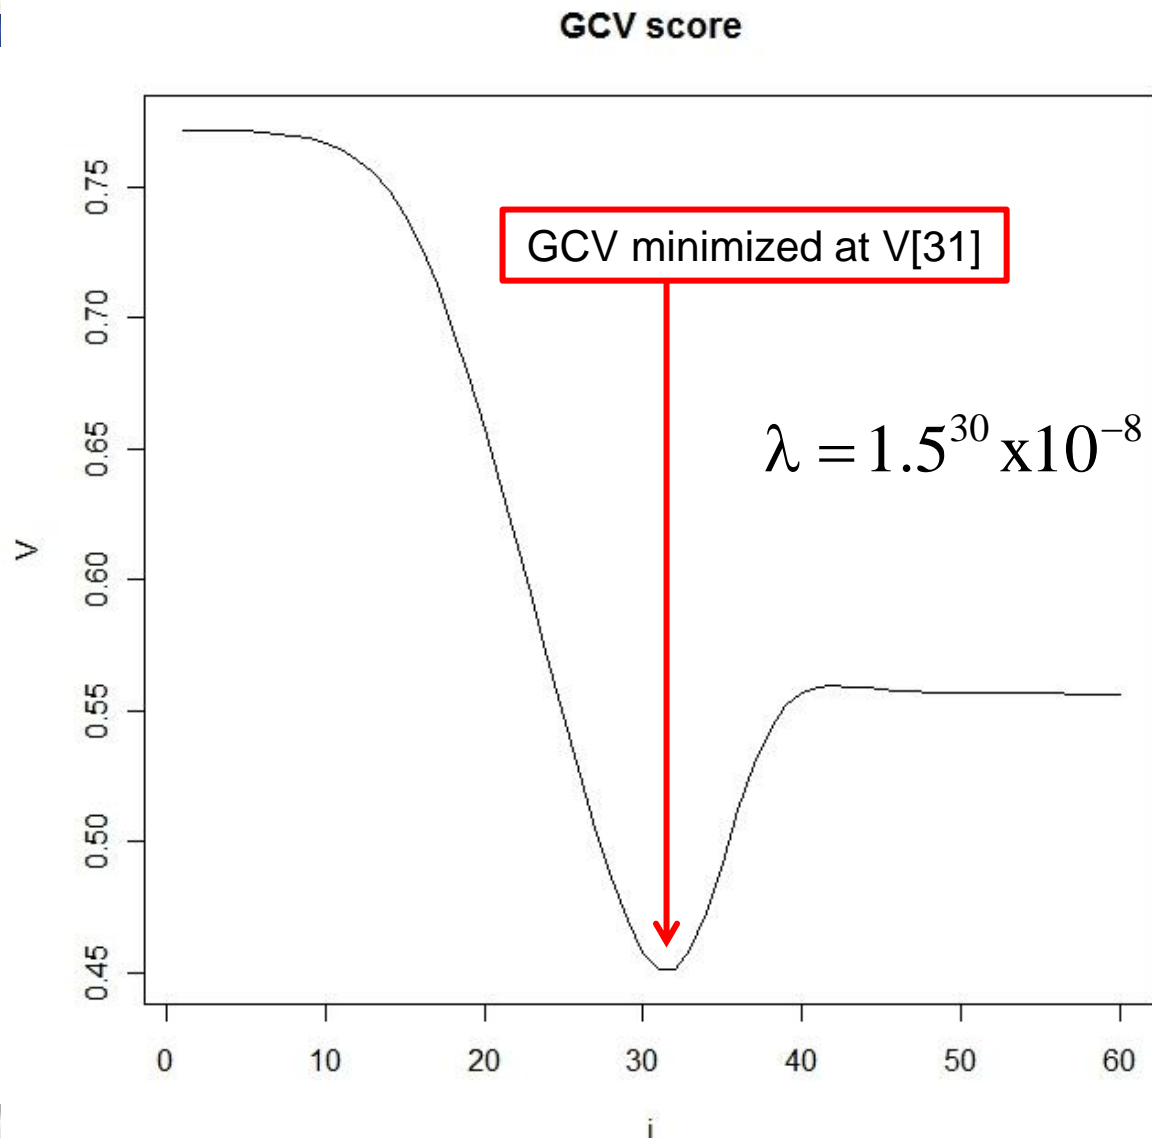
- Calculate **ordinary cross validation (OCV) score**: leave out each datum, fit the model, and calculate squared difference between missing datum and predicted value.
- We derive a proxy for M , V_0 , and choose λ to minimize V_0 .
- Leaving out one datum at a time and fitting the model to the remaining data points is **computationally inefficient**.
- It can be shown that you can use **mean weights** (instead of individual weights) to derive a similarly-purposed **generalized cross validation (GCV) score**.
 - GCV has **computational advantages** over OCV.
 - GCV also has advantages in terms of **invariance**.

R Code to Search for GCV Optimal Smoothing Parameter

- **Engine Wear:** Simple direct search for GCV optimal smoothing parameter:

```
lambda <- 1e-8; n <- length(wear); V <- rep(0,60)
for (i in 1:60)                # loop through smoothing parameters
{ mod<-prs.fit(wear,x,xk,lambda) # fit model, given lambda
  trA<-sum(influence(mod)$hat[1:n]) # find tr(A)
  rss<-sum((wear-fitted(mod)[1:n])^2) # residual sum of squares
  V[i]<-n*rss/(n-trA)^2             # obtain GCV score
  lambda<-lambda*1.5              # increase lambda
}
plot(1:60,V,type="l",main="GCV score",xlab="i") # plot score
```

GCV Function for Engine Wear



Lambda optimal to minimize GCV

X-axis reflects smoothing parameters on a log scale such that:

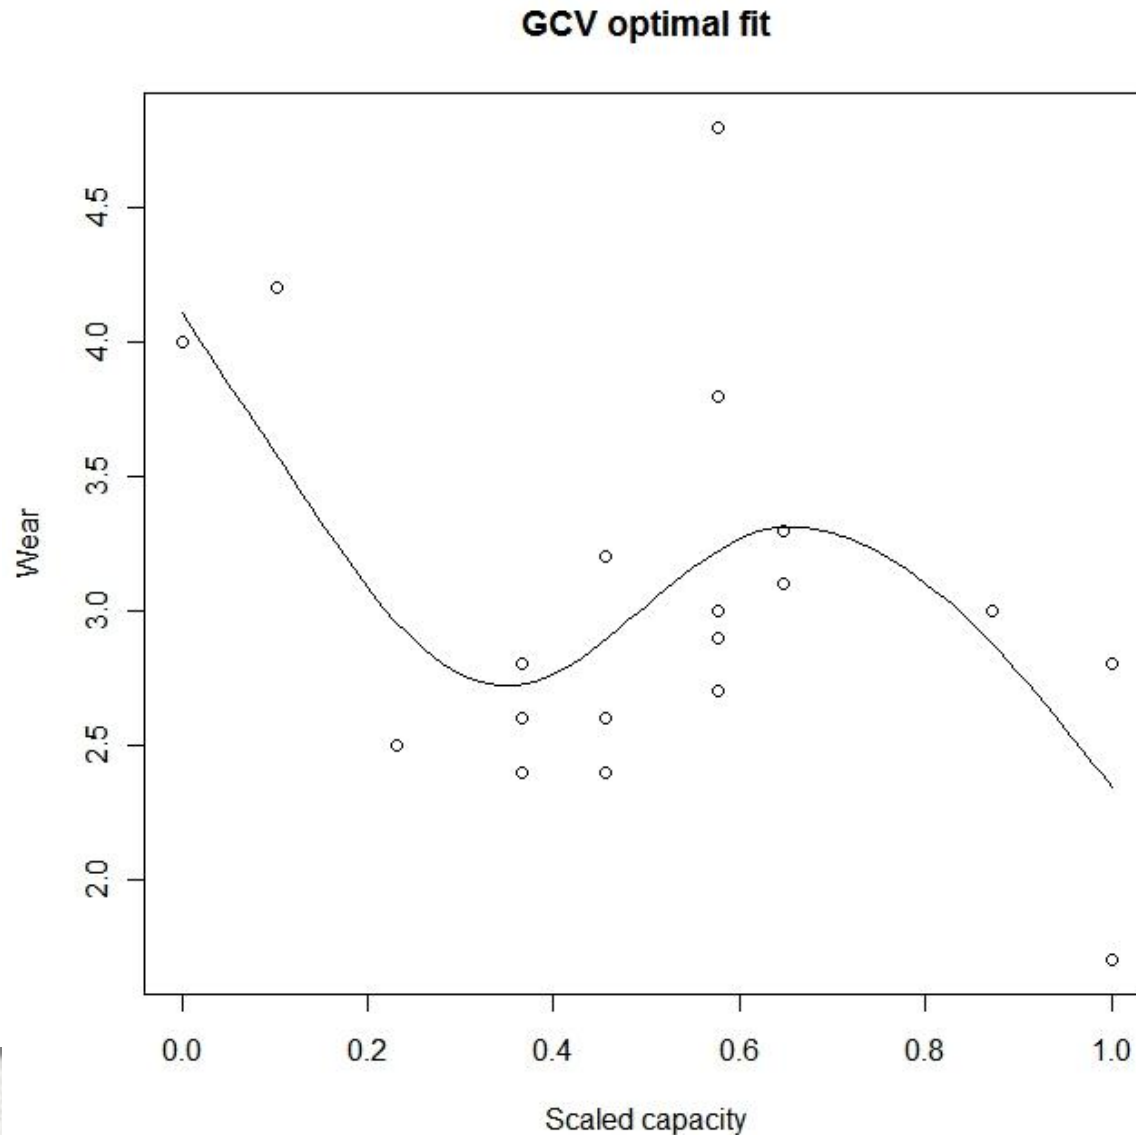
$$\lambda = 1.5^{i-1} \times 10^{-8}$$

R Code to Plot Optimal Model

- **Engine Wear:** Plot fitted model where GCV is minimized:

```
i<- (1:60) [V==min(V)] # extract index of min(V)
mod.3<-prs.fit(wear,x,xk,1.5^(i-1)*1e-8) # fit optimal model
Xp<-spl.X(xp,xk) # .... and plot it
plot(x,wear);lines(xp,Xp%*%coef(mod.3))
```

Engine Wear: Plot GCV-optimal fit $\lambda = 0.002$



Additive Models: Two Explanatory Variables

SIMON N. WOOD

- Suppose there are two explanatory variables, x and z , for a response variable, y , and so we use this simple additive model structure:

$$y_i = f_1(x_i) + f_2(z_i) + \varepsilon_i$$

Where f_1 and f_2 are smooth functions, and the ε_i are i.i.d. $N(0, \sigma^2)$ random normal variables and z_i and x_i lie in $[0, 1]$.

- First, the model above is a special restrictive case of $f(x, z)$
- The fact that the model contains more than one function introduces an identifiability problem:
 - For example, could simultaneously add any constant to f_1 and subtract it from f_2 without changing the model predictions.
 - Hence must impose identifiability constraints on the model before fitting.

Additive Models: Two Explanatory Variables



- Represent Each Smooth Function with a Penalized Regression Spline Basis:

$$f_1(x) = \delta_1 + x\delta_2 + \sum_{j=1}^{q_1-2} R(x, x_j^*) \delta_{j+2}$$

$$f_2(z) = \gamma_1 + z\gamma_2 + \sum_{j=1}^{q_2-2} R(z, z_j^*) \gamma_{j+2}$$

- Where δ_j and γ_j are the unknown parameters for f_1 and f_2 respectively, q_1 and q_2 are the number of unknown parameters for f_1 and f_2 , while x_j^* and z_j^* are the knot locations for the two functions.

Function to Set Up Simple Two Term Additive Model

```
## 3.3.1 Penalized regression spline representation of an additive model
am.setup <- function(x,z,q=10)
# Get X, S_1 and S_2 for a simple 2 term AM
{ # choose knots ...
  xk <- quantile(unique(x),1:(q-2)/(q-1)) # knot locations x
  zk <- quantile(unique(z),1:(q-2)/(q-1)) # knot locations z
  # get penalty matrices ...
  S <- list() # initializes S to be a list
  S[[1]] <- S[[2]] <- matrix(0,2*q-1,2*q-1) # initialize two components
  S[[1]][2:q,2:q] <- spl.S(xk)[-1,-1]
  S[[2]][(q+1):(2*q-1),(q+1):(2*q-1)] <- spl.S(zk)[-1,-1]
  # get model matrix ...
  n <- length(x)
  X <- matrix(1,n,2*q-1)
  X[,2:q] <- spl.X(x,xk)[,-1] # 1st smooth
  X[, (q+1):(2*q-1)] <- spl.X(z,zk)[,-1] # 2nd smooth
  list(X=X,S=S)
}
```

Function to Fit the Model and Calculate the GCV Score

```
fit.am <- function(y,X,S,sp)
# function to fit simple 2 term additive model
{ # get sqrt of total penalty matrix ...
  rS <- mat.sqrt(sp[1]*S[[1]]+sp[2]*S[[2]])
  q.tot <- ncol(X) # number of params
  n <- nrow(X) # number of data
  X1 <- rbind(X,rS) # augmented X
  y1 <- c(y,rep(0,q.tot)) # augment data
  b<-lm(y1~X1-1) # fit model
  trA<-sum(influence(b)$hat[1:n]) # tr(A)
  norm<-sum((y-fitted(b)[1:n])^2) # RSS
  list(model=b,gcv=norm*n/(n-trA)^2,sp=sp)
}
```

Estimate an Additive Model to Fit Data for 31 Felled Cherry Trees

- We will use this routine to estimate an additive model for the data in data frame: `trees`.

- `Volume`, `Girth` and `Height` for 31 felled cherry trees.
- We want to predict `Volume` using the model:

$$\text{Volume} = f_1(\text{Girth}) + f_2(\text{Height}) + \varepsilon_i$$

- Given the simple smoothers we are using, we must rescale the predictors onto $[0,1]$:

```
data(trees) # get tree data
rg <- range(trees$Girth) # rg has low, high of girth
trees$Girth <- (trees$Girth - rg[1]) / (rg[2] - rg[1])
rh <- range(trees$Height)
trees$Height <- (trees$Height - rh[1]) / (rh[2] - rh[1])
```

- Then we obtain the model matrix and penalty matrices: 36
- ```
am0 <- am.setup(trees$Girth, trees$Height)
```



# Grid Search to Find Model Fit

## Minimizing GCV Score

- Now we perform a grid search to find the model fit that approximately minimizes the GCV score:

```
sp <-c(0,0) # initialize smoothing parameter (s.p.) array
for (i in 1:30) for (j in 1:30) # loop over s.p. grid
{ sp[1]<-1e-5*2^(i-1);sp[2]<-1e-5*2^(j-1) # s.p.s
 b<-fit.am(trees$Volume,am0$X,am0$S,sp) # fit using s.p.s
 if (i+j==2) best<-b else # store 1st model
 if (b$gcv<best$gcv) best<-b # store best model
}
```

best\$sp # GCV best smoothing parameters found

[1] 0.01024 5368.70912

Height has **high** smoothing parameter (straight line)

Girth has **low** smoothing parameter (allows curvature)

# Grid Search to Find Model Fit

## Minimizing GCV Score

- So smooth of girth has a fairly low smoothing parameter (0.01), allowing  $f_1$  some curvature, whereas  $f_2$  has a very high smoothing parameter (5368.71).
- We obtain values of the smooths at predictor variable values by (1) zeroing all model coefficients (except those corresponding to the term of interest, and (2) using `predict()`: Next slide . . .

# Plot: 1. Actual Volume versus Fitted; 2. $f_1$ Predicted by Scaled Girth

- Determine values of smooths at predictor variable values by zeroing all model coefficients (except predicted value):

```
plot fitted volume against actual volume ...
```

```
plot(trees$Volume,fitted(best$model)[1:31],
 xlab="Fitted Volume",ylab="Actual Volume")
```

```
evaluate and plot f_1 against Girth ...
```

```
b <-best$model
```

```
b$coefficients[1]<-0 # zero the intercept
```

```
b$coefficients[11:19]<-0 # zero the second smooth coefs
```

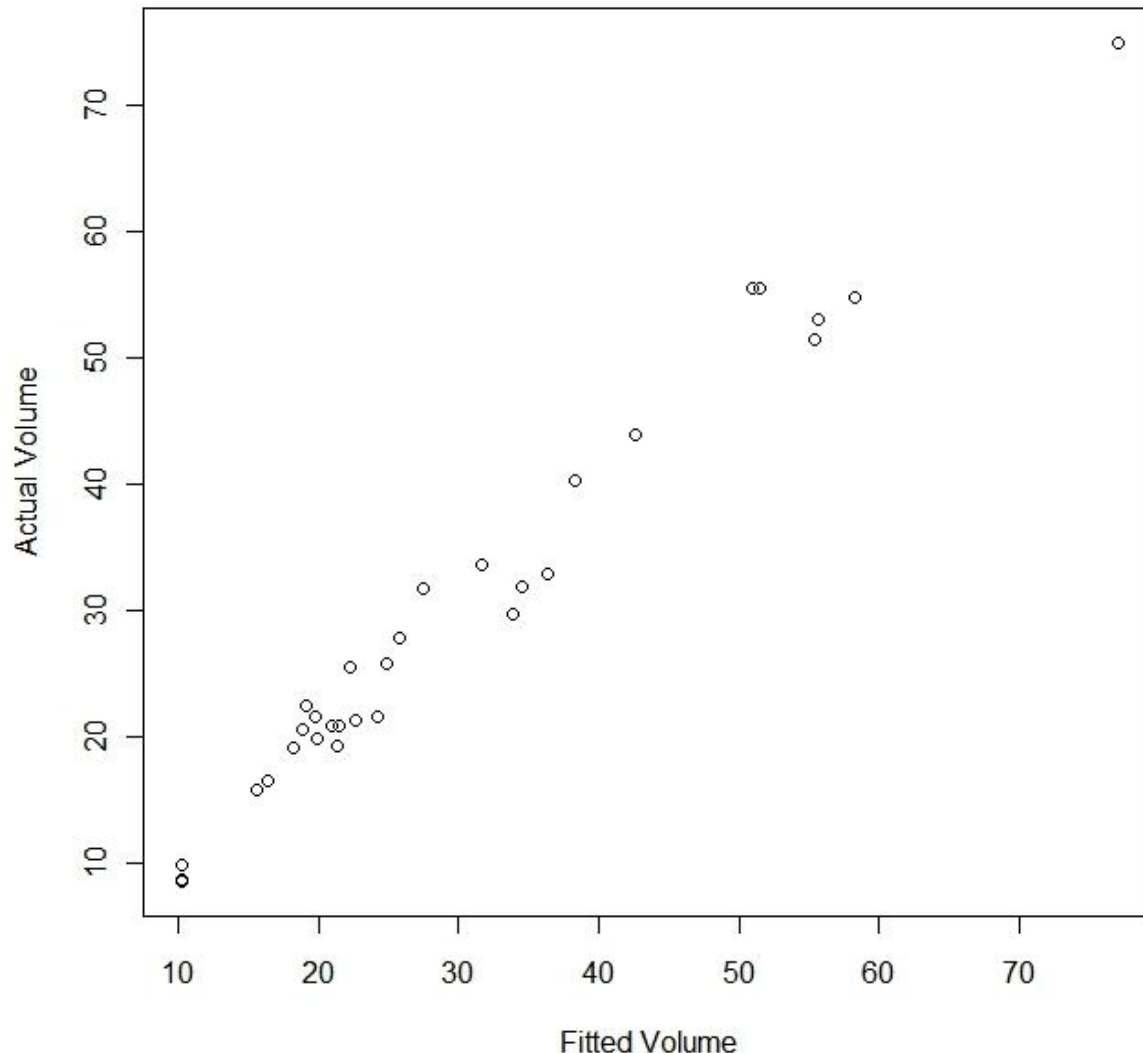
```
f0<-predict(b) # predict f_1 only, at data values
```

```
plot(trees$Girth,f0[1:31],xlab="Scaled Girth",
 ylab=expression(hat(f[1])))
```

# Plot Actual Volume versus Fitted

```
plot(trees$Volume,fitted(best$model)[1:31],main= "ACTUAL VOLUME
BY FITTED VOLUME", xlab="Fitted Volume", ylab="Actual Volume")
```

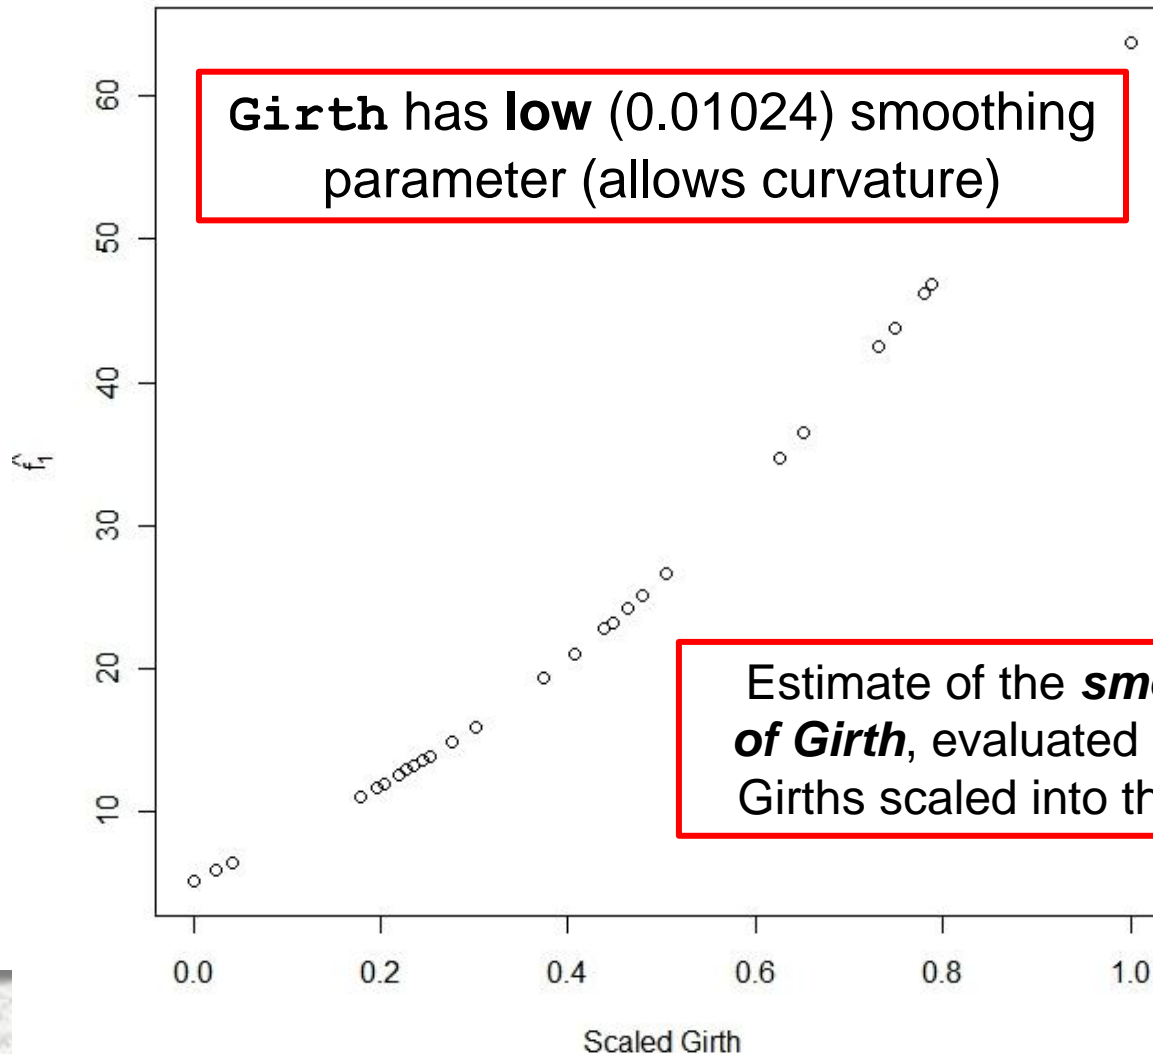
ACTUAL VOLUME BY FITTED VOLUME



# Plot $\hat{f}_1$ Predicted by Scaled Girth

```
plot(trees$Girth,f0[1:31],main= "f_1 Predicted by Scaled
Girth",xlab="Scaled Girth", ylab=expression(hat(f[1])))
```

$\hat{f}_1$  Predicted by Scaled Girth



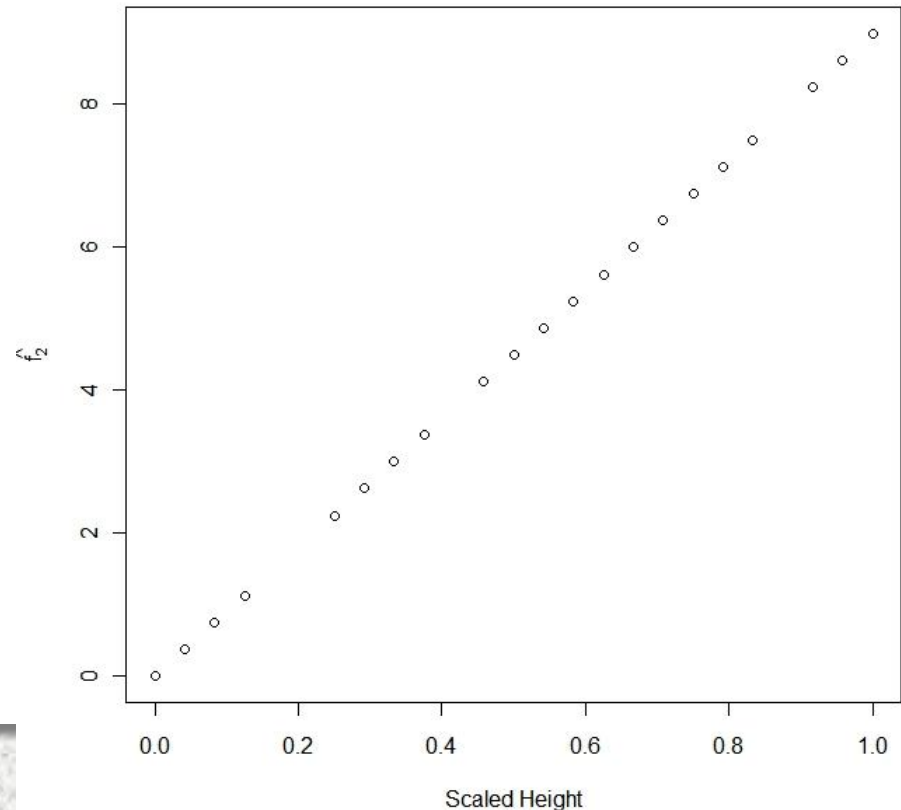
Estimate of the ***smooth function of Girth***, evaluated at the observed Girths scaled into the  $[0,1]$  interval.

# Plot $f_2$ Predicted by Scaled Height

```
evaluate and plot f_2 against Height ...
b <- best$model
b$coefficients[1] <- 0 # zero the intercept
b$coefficients[2:10] <- 0 # zero the first smooth coefs
f0 <- predict(b) # predict f_2 only, at data values
plot(trees$Height, f0[1:31], main = "f_2 Predicted by Scaled
 Height", xlab = "Scaled Height", ylab = expression(hat(f[2])))
f_2 Predicted by Scaled Height
```

Estimate of the ***smooth function of Height***, evaluated at the observed Heights scaled into the  $[0,1]$  interval.

**Height** has **high** (5368) smoothing parameter (disallows curvature).





# Generalized Additive Models

- **Generalized Additive Models (GAMs)** follow from **Additive Models**, as **GLMs** follow from **Linear Models**.
- **Linear Predictor** predicts some known *smooth monotonic function* of the **Expected Value** of the *response*, which
  - Follows any exponential family distribution; or
  - Has a known mean variance relationship.
- Consider **trees** data again using GAM of form:

$$\log\{E(\text{Volume}_i)\} = f_1(\text{Girth}_i) + f_2(\text{Height}_i), \text{ Volume} \sim \text{Gamma}$$

# tree GAM: What Needs To Be Done?

- Go back and look at `fit.am()` on slide 35:

```
fit.am <- function(y,X,S,sp)
function to fit simple 2 term additive model
{ # get sqrt of total penalty matrix ...
 rS <- mat.sqrt(sp[1]*S[[1]]+sp[2]*S[[2]])
 q.tot <- ncol(X) # number of params
 n <- nrow(X) # number of data
 X1 <- rbind(X,rS) # augmented X
 y1 <- c(y,rep(0,q.tot)) # augment data
 b<-lm(y1~X1-1) b<-glm(y1~X1-1)? NO! # fit model
 trA<-sum(influence(b)$hat[1:n]) # tr(A)
 norm<-sum((y-fitted(b)[1:n])^2) # RSS
 list(model=b,gcv=norm*n/(n-trA)^2,sp=sp)
}
```

Additive model was estimated by **penalized least squares**.  
GAM will be fitted by **penalized likelihood maximization**.  
Using **penalized iteratively re-weighted least squares (P-IRLS)**

# tree GAM: What Needs To Be Done?

- Modify `fit.am()` to `fit.gamG()` below:

```
fit.gamG <- function(y,X,S,sp)
function to fit simple 2 term generalized additive model
Gamma errors and log link
{ # get sqrt of combined penalty matrix ...
 rS <- mat.sqrt(sp[1]*S[[1]]+sp[2]*S[[2]])
 q <- ncol(X) # number of params
 n <- nrow(X) # number of data
 X1 <- rbind(X,rS) # augmented model matrix
 eta <- log(y) # initialize linear predictor
 norm <- 0;old.norm <- 1 # initialize convergence control
 while (abs(norm-old.norm)>1e-4*norm) # repeat un-converged
 { mu <- exp(eta) # fitted values
 z <- (y-mu)/mu + eta # pseudodata (recall w_i=1, here)
 z[(n+1):(n+q)] <- 0 # augmented pseudodata
 m <- lm(z~X1-1) # fit penalized working model
 b <- m$coefficients # current parameter estimates
 eta <- (X1%*%b)[1:n] # 'linear predictor'
 trA <- sum(influence(m)$hat[1:n]) # tr(A)
 old.norm <- norm # store for convergence test
 norm <- sum((z-fitted(m))[1:n]^2) # RSS of working model
 }
 list(model=m,gcv=norm*n/(n-trA)^2,sp=sp)
}
```

# tree GAM: What Needs To Be Done?

- To find GCV optimum fit, simply **replace `fit.am()` function with `fit.gamG()` function** in the smoothing parameter grid search loop (slide 37) and repeated below:

```
to find model fit that approximately minimizes GCV score
sp <- c(0,0) # initialize smoothing parameter (s.p.) array
for (i in 1:30) for (j in 1:30) # loop over s.p. grid
{ sp[1]<-1e-5*2^(i-1);sp[2]<-1e-5*2^(j-1) # s.p.s
 b<-fit.am(trees$Volume,am0$X,am0$S,sp) # fit using s.p.s
 if (i+j==2) best<-b else # store 1st model
 if (b$gcv<best$gcv) best<-b # store best model
}
best$sp # GCV best smoothing parameters found
```

**LET'S TRY IT !!!**

# tree GAM: *We Try It!*

```
to find model fit that approximately minimizes GCV score for
a GAM, we have
replaced fit.am with fit.gamG in the smooting parameter grid
search loop below:
sp <- c(0,0) # initialize smoothing parameter (s.p.) array
for (i in 1:30) for (j in 1:30) # loop over s.p. grid
{ sp[1]<-1e-5*2^(i-1);sp[2]<-1e-5*2^(j-1) # s.p.s
 b<-fit.gamG(trees$Volume,am0$X,am0$S,sp) # fit using s.p.s
 if (i+j==2) best<-b else # store 1st model
 if (b$gcv<best$gcv) best<-b # store best model
}
best$sp # GCV best smoothing parameters found
[1] 0.01024 5368.70912
```



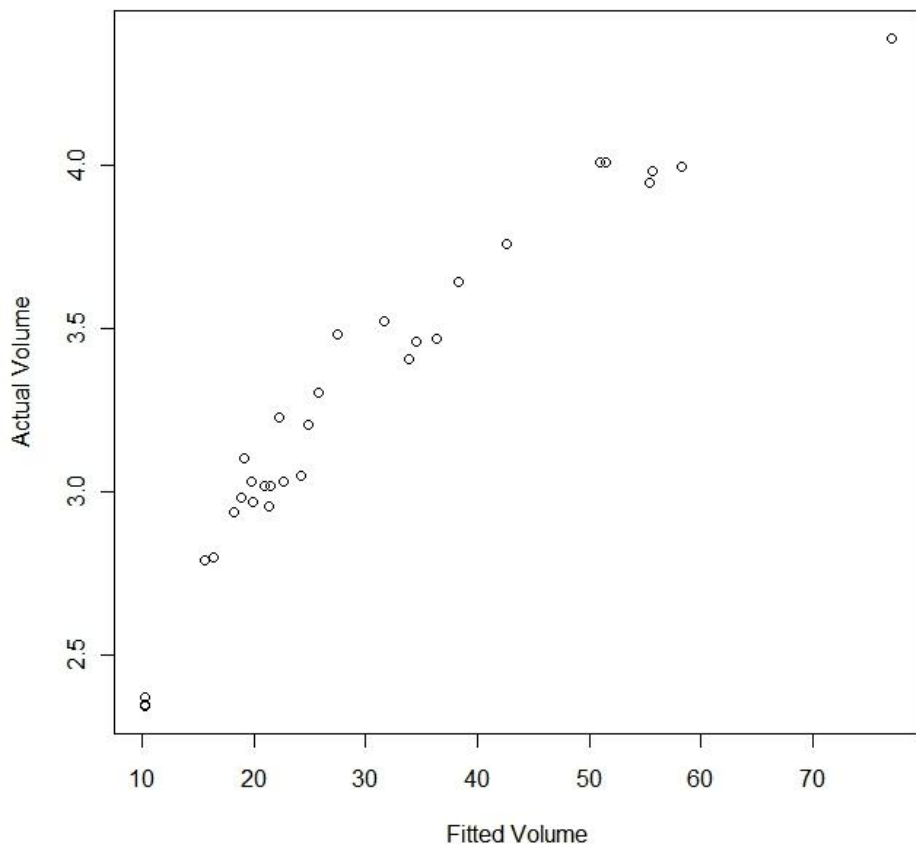
We get the same optimum smoothing GCV scores

# Plot Actual Volume versus Fitted

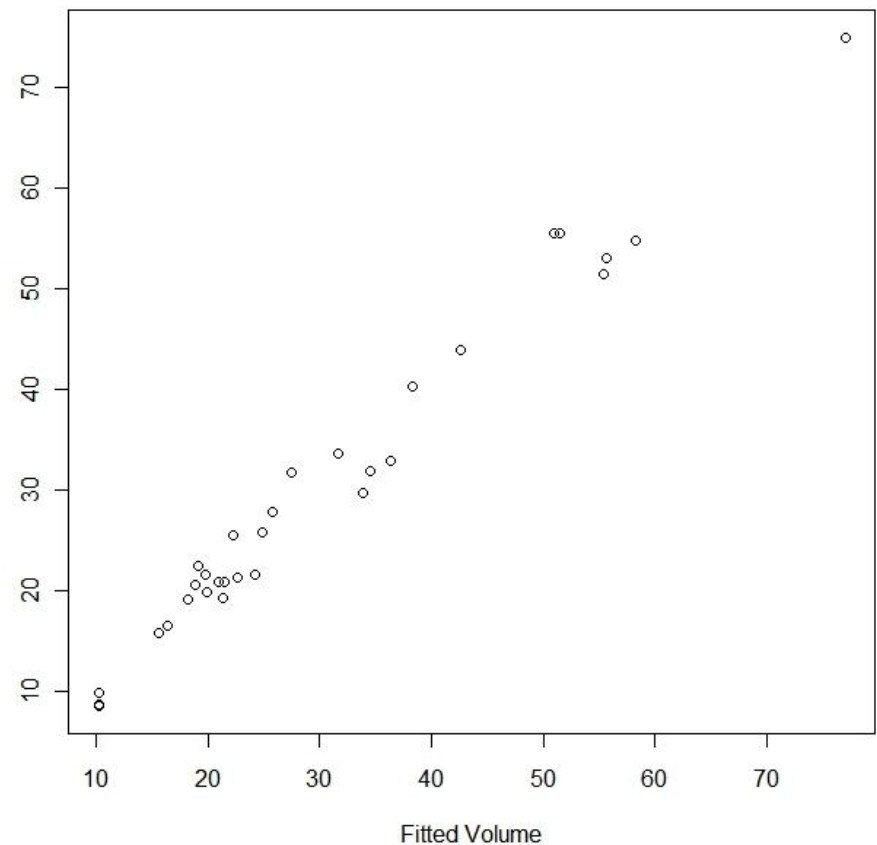
```
plot(trees$Volume,fitted(best$model)[1:31],main= "ACTUAL VOLUME
BY FITTED VOLUME", xlab="Fitted Volume", ylab="Actual Volume")
```

They are the same, but they should be !

ACTUAL VOLUME BY FITTED VOLUME: GAM



ACTUAL VOLUME BY FITTED VOLUME





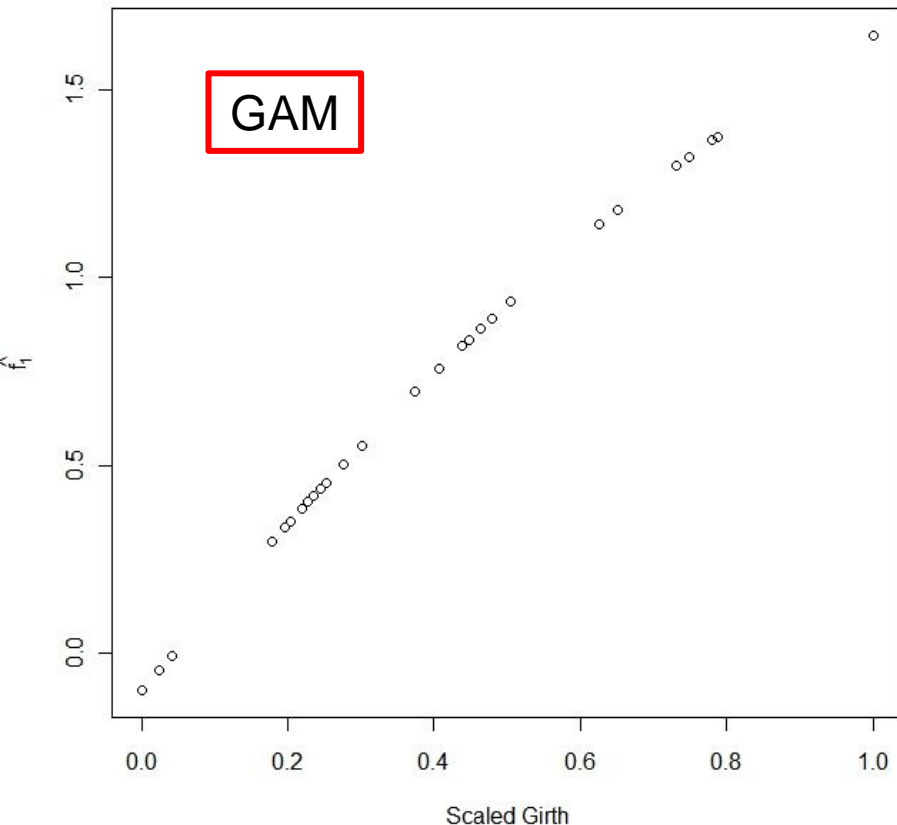
# Plot $f_1$ Predicted by Scaled Girth

```
plot(trees$Girth,f0[1:31],main= "f_1 Predicted by Scaled
Girth",xlab="Scaled Girth", ylab=expression(hat(f[1])))
```

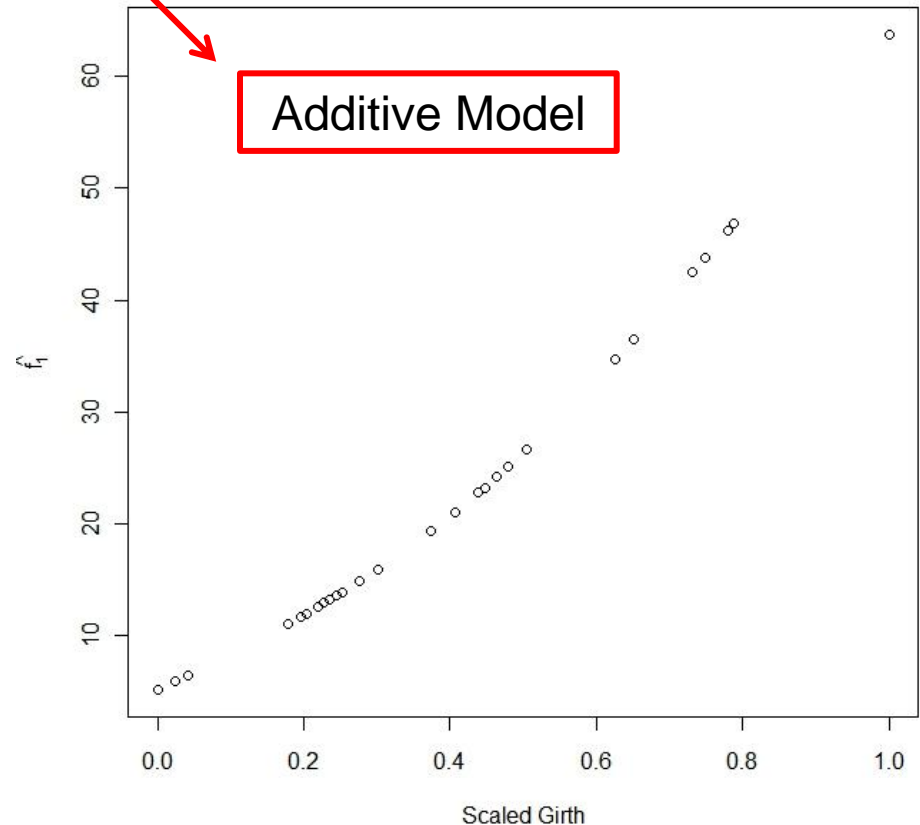
**Girth** has **low** (0.01024) smoothing parameter (allows curvature)

Estimate of the ***smooth function of Girth***, evaluated at the observed Girths scaled into the  $[0,1]$  interval.

f\_1 Predicted by Scaled Girth: GAM



f\_1 Predicted by Scaled Girth



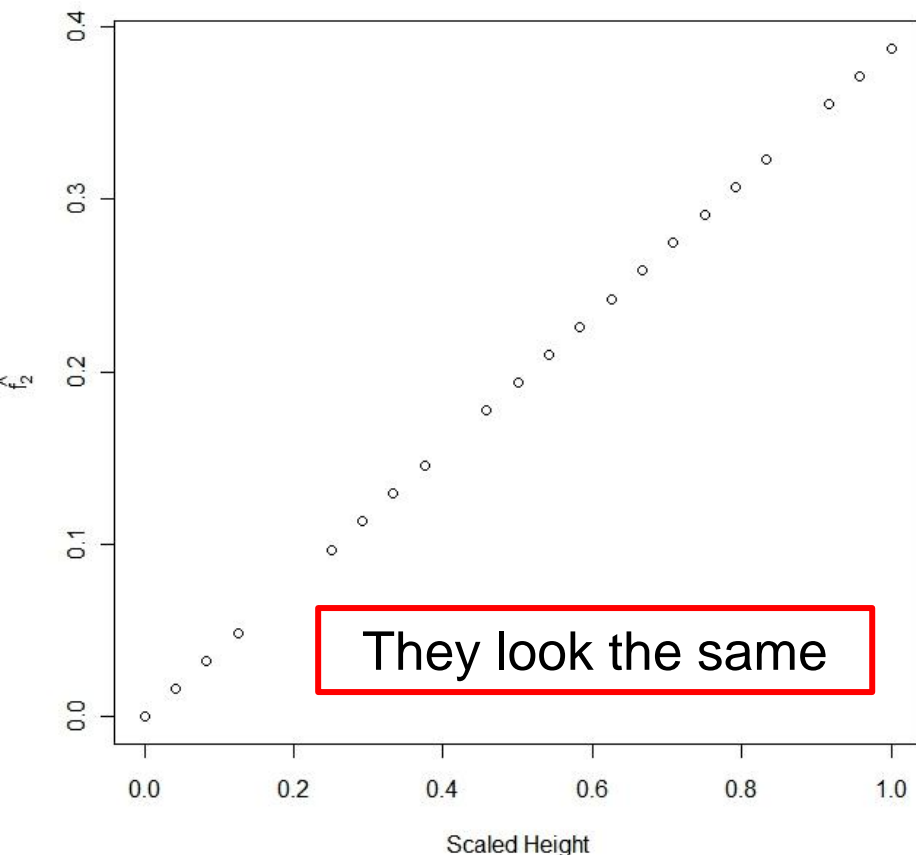
# Plot $f_2$ Predicted by Scaled Height

```
plot(trees$Height,f0[1:31],main="f_2 Predicted by Scaled Height:
GAM",xlab="Scaled Height",ylab=expression(hat(f[2])))
```

**Height** has **high** (5368) smoothing parameter (disallows curvature).

Estimate of the ***smooth function of Height***, evaluated at the observed Heights scaled into the [0,1] interval.

f\_2 Predicted by Scaled Height: GAM



f\_2 Predicted by Scaled Height

